

Context-Window Degradation and Its Impact on Long-Document Research Synthesis

Abstract

Large language models (LLMs) have rapidly developed the ability to process increasingly long textual contexts, creating new possibilities for academic research, document analysis, literature review, and evidence synthesis. However, an important distinction exists between the **maximum context window** advertised by a model and its **effective ability to use information reliably throughout that context**. Research has demonstrated that model performance may decline as input length increases, particularly when relevant information is located in the middle of a long context or when a task requires information to be retrieved and integrated across multiple documents. This phenomenon can be described broadly as **context-window degradation**.

The problem is particularly important for long-document research synthesis. Academic synthesis requires more than retrieving an isolated fact: a system must identify relevant evidence, preserve relationships between claims, compare sources, resolve contradictions, and construct a coherent interpretation. Long contexts can therefore create problems involving information retrieval, attention allocation, evidence coverage, reasoning, and hallucination. Studies such as Liu et al. (2024), Bai et al. (2024), and Hsieh et al. (2024) show that long-context performance can deteriorate even when models appear capable of processing very large inputs. At the same time, approaches such as retrieval augmentation, context compression, hierarchical processing, and structured decomposition offer possible solutions.

This paper examines the causes and consequences of context-window degradation for long-document research synthesis and discusses strategies for improving reliability.

Keywords: large language models, context window, long-context learning, document synthesis, retrieval-augmented generation, hallucination, information retrieval, academic research

1. Introduction

Large language models have become increasingly useful for processing documents, answering questions, summarizing research, and supporting literature analysis. One of the major developments in this area has been the expansion of model

context windows. Earlier language models were restricted to relatively short sequences, whereas modern systems can process tens or hundreds of thousands of tokens, with some systems supporting substantially larger contexts.

The expansion of context windows creates an apparent solution to an important problem: if a model can read an entire report, book, collection of research papers, or legal document at once, it should theoretically be able to synthesize information from the complete source material. However, empirical research suggests that **context capacity and context utilization are not equivalent**.

Liu et al. (2024), for example, examined how language models use information distributed throughout long contexts. Their experiments demonstrated that models can perform substantially worse when relevant information is positioned in the middle of a long input. This phenomenon has become widely known as **“lost in the middle.”** The finding is important because it demonstrates that simply increasing the number of tokens a model can accept does not guarantee uniform access to the information contained within those tokens.

This problem becomes even more significant in research synthesis. A research synthesis task may require a model to identify evidence from several parts of a document, compare findings across different studies, recognize methodological differences, and combine evidence into a defensible conclusion. Consequently, long-document synthesis places demands on models that are substantially more complex than simple fact retrieval.

The central argument of this assignment is that **context-window degradation can reduce the reliability, completeness, and evidential grounding of long-document research synthesis even when the model technically accepts the entire document as input**.

2. Understanding Context-Window Degradation

2.1 What is a context window?

A context window is the amount of tokenized information that a language model can consider during a particular interaction or inference process. It can include the user’s question, instructions, previous conversation, retrieved documents, and other textual material.

A larger context window provides greater capacity for processing long documents. However, it should not be interpreted as equivalent to perfect comprehension of the entire context.

The distinction can be represented as:

Maximum context length ≠ effective context utilization

A model may technically accept a 100,000-token document while still demonstrating weaker retrieval or reasoning performance as the amount of contextual information increases.

2.2 The “lost in the middle” phenomenon

One of the most influential findings concerning long-context degradation comes from Liu et al. (2024). Their experiments investigated multi-document question answering and key-value retrieval while changing the position of relevant information in the context.

The researchers found that performance was often strongest when relevant information occurred near the beginning or end of the input and substantially weaker when the information appeared in the middle.

This finding has direct implications for academic research. A relevant finding in the middle of a 200-page report may be less reliably utilized than an identical finding placed near the beginning or end of the supplied context.

Therefore, long-document processing is not necessarily a uniform process in which every token receives equal functional importance.

3. Why Long Documents Create Problems for Research Synthesis

Research synthesis differs from ordinary summarization because it requires integration of evidence.

For example, consider a researcher asking an LLM to analyze ten papers concerning artificial intelligence in education. A reliable synthesis may require the model to:

1. identify the research question of each study;
2. distinguish empirical findings from theoretical claims;
3. identify the methodology used;
4. compare sample sizes and populations;
5. identify conflicting results;
6. assess limitations;
7. connect findings across papers; and
8. develop an evidence-supported overall conclusion.

When the information is distributed across a very long context, each of these operations becomes more difficult.

3.1 Retrieval degradation

The first challenge is finding the relevant evidence.

Long contexts contain a high proportion of information that may be irrelevant to a particular question. As context size increases, the model must distinguish relevant information from increasingly large quantities of distractor material.

Hsieh et al. (2024), through the RULER benchmark, demonstrated that models that perform extremely well on simple needle-in-a-haystack retrieval can show significant performance degradation as context length and task complexity increase.

This suggests that simple retrieval tests may overestimate genuine long-context understanding.

3.2 Evidence coverage

A synthesis can fail even when every individual statement generated by the model appears plausible.

Suppose an academic report contains 15 relevant findings, but the model identifies only six of them. The resulting synthesis may be fluent and factually reasonable while nevertheless being incomplete.

This creates an important distinction:

Answer correctness $\neq$ evidence completeness

For academic research, omission can be as important as factual error.

3.3 Cross-document reasoning

Long-document synthesis frequently requires combining evidence from different sources.

For example:

Paper A reports a positive effect.

Paper B reports no significant effect.

Paper C suggests that the effect depends on age.

A useful synthesis must not simply list these findings. It must recognize the relationship between them and potentially identify methodological or contextual reasons for disagreement.

This makes research synthesis a **multi-hop reasoning problem**, rather than simple information retrieval.

4. Evidence from Long-Context Benchmarks

Several benchmarks have been developed to measure long-context capabilities.

4.1 LongBench

Bai et al. (2024) introduced **LongBench**, a bilingual, multitask benchmark for evaluating long-context understanding. It includes tasks involving single-document question answering, multi-document question answering, summarization, few-shot learning, synthetic tasks, and code completion.

The importance of LongBench is that it evaluates multiple forms of long-context understanding rather than relying exclusively on simple retrieval.

The study found that even capable models could struggle with longer contexts. It also showed that retrieval and context-compression methods can improve performance for models with weaker long-context capabilities.

4.2 RULER

Hsieh et al. (2024) introduced **RULER**, designed to provide a more comprehensive evaluation of long-context capabilities.

The authors evaluated 17 long-context language models across 13 tasks. Importantly, they found that models could achieve very high performance on traditional needle-in-a-haystack tests while exhibiting substantial performance degradation when context length and task complexity increased.

This finding is particularly important for research synthesis because real academic tasks are rarely equivalent to finding one isolated sentence. Research requires aggregation, tracing relationships, comparison, and reasoning.

4.3 LooGLE and related evaluations

Other benchmarks have similarly emphasized that real-world long-document tasks involve more complex dependencies than artificial retrieval tests. Recent evaluations increasingly distinguish between **memorizing or retrieving information** and **reasoning over relationships distributed throughout long documents**.

Consequently, future evaluation of research assistants should measure not only whether a model can locate a fact but also whether it can provide comprehensive, source-grounded synthesis.

5. Mechanisms Behind Context-Window Degradation

Several mechanisms may contribute to degradation.

5.1 Attention allocation

Transformer-based models use attention mechanisms to process relationships among tokens. As sequence length increases, the model must process a much larger number of potential relationships.

The resulting challenge is not necessarily that information disappears from the input. Rather, the model may fail to allocate sufficient effective attention to the information required by the task.

5.2 Positional effects

Where information appears within a context can influence model performance. The “lost in the middle” findings provide evidence that position matters.

This means that two contexts containing exactly the same information may produce different results if the relevant information is positioned differently.

5.3 Distractor information

Long documents contain redundant, irrelevant, and sometimes contradictory information.

A model must distinguish:

- relevant evidence,
- background information,
- methodological details,
- examples,
- repeated claims,
- contradictory findings, and
- irrelevant material.

As the amount of distractor information increases, the difficulty of selecting appropriate evidence also increases.

5.4 Multi-step reasoning

Long-document synthesis often requires chains of reasoning. A model may need to retrieve one finding, connect it to a methodological limitation elsewhere, compare it with another paper, and then derive a conclusion.

Failure at any stage can affect the final synthesis.

Thus:

Long-context synthesis $\approx$ retrieval + comparison + reasoning + evidence integration

Increasing context length can increase the difficulty of all four components.

6. Impact on Academic Research

Context-window degradation can have several consequences for researchers.

6.1 Incomplete literature reviews

If a model fails to retrieve relevant information from portions of a long document or collection of papers, its literature review may omit important evidence.

This is especially problematic when the omitted evidence contradicts the dominant conclusion.

6.2 Misrepresentation of scholarly consensus

Suppose eight papers support a conclusion while four papers report contradictory results. If the model disproportionately retrieves the first group, it may incorrectly describe the literature as more consistent than it actually is.

Therefore, context degradation can contribute to **false impressions of consensus**.

6.3 Loss of methodological nuance

Academic findings cannot always be interpreted independently of methodology.

For example, two studies might reach different conclusions because they use different:

- populations,
- sample sizes,
- measurement instruments,
- experimental designs,
- time periods, or
- statistical methods.

If methodological details are overlooked, synthesis may become misleading.

6.4 Hallucination and unsupported inference

When relevant evidence is not effectively retrieved, a model may generate a plausible connection based on patterns learned during training rather than evidence present in the supplied documents.

This creates a grounding problem.

A useful principle for academic AI systems is:

The longer the context, the more important explicit evidence tracking becomes.

7. Context Compression as a Possible Solution

One response to context degradation is to reduce the amount of information supplied to the model while preserving the information most relevant to the task.

LLMLingua and related context-compression approaches demonstrate this direction. Instead of passing an entire lengthy prompt unchanged, a system can identify and retain information considered important for the downstream task.

Context compression can provide two benefits:

1. reduced computational cost; and
2. reduced exposure to irrelevant information.

However, compression introduces its own risk: **the compression process may remove information that later becomes important.**

Therefore, the objective should not simply be:

“Make the context shorter.”

Instead, the objective should be:

“Reduce irrelevant information while preserving evidence required for the research task.”

8. Retrieval-Augmented Generation

Retrieval-augmented generation (RAG) provides another important strategy.

Instead of giving the model an entire collection of documents, a RAG system retrieves passages that are likely to be relevant to the user's question and provides those passages to the language model.

A simplified workflow is:

Research question → Document retrieval → Relevant passages → LLM reasoning → Answer with evidence

This can reduce the amount of irrelevant material that the model must process.

However, retrieval itself creates a potential failure point. If the retriever fails to select an important passage, the language model may never receive the evidence required for a correct synthesis.

Thus, RAG changes the problem from:

“Can the model use everything?”

to:

“Can the system retrieve everything important and then reason correctly over it?”

For research synthesis, this distinction is critical.

9. Hierarchical and Iterative Synthesis

A further strategy is to process documents hierarchically rather than placing everything into one context.

For example:

Level 1: Summarize individual sections.

↓

Level 2: Compare section summaries.

↓

Level 3: Identify cross-section themes.

↓

Level 4: Compare themes across documents.

↓

Level 5: Produce the final synthesis.

This approach reduces the amount of information that must be handled simultaneously.

However, hierarchical synthesis introduces a potential **error propagation problem**. If an important fact is omitted during an early summary, it may be unavailable to later stages.

Consequently, hierarchical systems should ideally preserve citations, source locations, and evidence trails throughout the process.

10. Evidence-Centered Research Synthesis

For academic applications, a particularly promising approach is to treat synthesis as an evidence-management problem rather than simply a generation problem.

A robust system could maintain an evidence table containing:

Source	Claim	Evidence	Location	Method	Limitations	Contradictory evidence
Study A	X improves performance	Experimental result	p. 12	RCT	Small sample	Study C
Study B	No significant effect	Statistical result	p. 8	Survey	Self-report	Study A
Study C	Effect depends on age	Subgroup analysis	p. 17	Longitudinal	Attrition	—

Such a structure can help prevent the model from treating a long document as an undifferentiated block of text.

It also makes the final synthesis easier to verify.

11. Implications for Academic Researchers

The literature suggests that researchers should avoid assuming that an LLM has comprehensively understood a document merely because the document fits within the model's context window.

Researchers should instead:

1. divide very long documents into meaningful sections;
2. retrieve evidence relevant to specific research questions;
3. preserve source references;
4. compare retrieved evidence with the original documents;
5. explicitly search for contradictory findings;
6. distinguish retrieval from synthesis;
7. verify important claims manually;
8. evaluate evidence coverage rather than only answer fluency.

This is especially important for systematic reviews and assignments where citations must be accurate.

12. Research Gaps

Although significant progress has been made, several research gaps remain.

12.1 Effective context length

A major unresolved question is how to define the **effective context length** of a model.

The advertised context window is a technical limit, whereas effective context length should represent the amount of information that a model can reliably retrieve and reason over.

12.2 Measuring synthesis quality

Many benchmarks focus on retrieval or question answering. Academic research synthesis requires additional measures such as:

- evidence coverage,
- citation accuracy,
- contradiction detection,
- methodological interpretation,
- claim attribution, and
- completeness.

12.3 Long-context hallucination

More research is needed to determine how context length interacts with hallucination. A model may have more evidence available but still generate unsupported claims if it fails to retrieve or correctly integrate the evidence.

12.4 Human-AI collaboration

Future research should investigate how researchers can work with LLMs most effectively. Instead of asking a model to independently summarize an entire literature corpus, interactive systems may allow researchers to inspect retrieved evidence, identify missing sources, and request targeted synthesis.

12.5 Domain-specific effects

Context degradation may differ across disciplines. Legal research, medicine, engineering, social science, and literature studies have different requirements for evidence and reasoning. More domain-specific evaluations are therefore necessary.

13. Conceptual Framework

The overall relationship can be represented as follows:

Increasing document length

↓

More information + more distractors

↓

Greater retrieval and attention demands

↓

Potential context-window degradation

↓

Reduced evidence retrieval

↓

Incomplete or distorted synthesis

↓

Potentially weaker conclusions

This relationship can be moderated by:

Retrieval augmentation + context compression + hierarchical processing + evidence tracking + human verification

These strategies can potentially improve the reliability of long-document research synthesis.

14. Conclusion

Context-window expansion is an important development in large language models, but the ability to accept more tokens should not be confused with the ability to reliably understand and synthesize all information contained in those tokens.

Research on “lost in the middle,” LongBench, and RULER demonstrates that long-context performance can decline as context length and task complexity increase. These findings are especially relevant to academic research because research

synthesis requires more than retrieving isolated facts. It requires comprehensive evidence retrieval, comparison of sources, recognition of contradictions, methodological interpretation, and multi-step reasoning.

The central challenge is therefore not simply **how much information an LLM can read**, but **how reliably it can identify, retain, connect, and use the information that matters**.

Approaches such as retrieval-augmented generation, context compression, hierarchical synthesis, and explicit evidence tracking provide promising directions. However, these techniques also introduce new risks, particularly retrieval failures and information loss during compression or summarization.

For academic research, the most reliable future systems are therefore likely to combine long-context models with structured retrieval, evidence preservation, iterative reasoning, and human verification. Context-window degradation should be treated not merely as a technical limitation of language models but as an important methodological issue affecting the reliability of AI-assisted research synthesis.

References

- Bai, Y., Lv, X., Zhang, J., Lyu, H., Tang, J., Huang, Z., Du, Z., Liu, X., Zeng, A., Hou, L., Dong, Y., Tang, J., & Li, J. (2024). *LongBench: A bilingual, multitask benchmark for long context understanding*. Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL).
<https://aclanthology.org/2024.acl-long.172/>
- Hsieh, C.-Y., Sun, S., Krizan, S., Acharya, S., Agarwal, D., Cai, J., et al. (2024). *RULER: What's the real context size of your long-context language models?* arXiv.
<https://arxiv.org/abs/2404.06654>
- Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., & Liang, P. (2024). *Lost in the middle: How language models use long contexts*. Transactions of the Association for Computational Linguistics, 12, 157–173.
<https://aclanthology.org/2024.tacl-1.9/>
- Zhang, Z., Chen, R., Liu, S., Yao, Z., et al. (2024). *Found in the middle: How language models use long contexts better via plug-and-play positional encoding*. Advances in Neural Information Processing Systems.
https://proceedings.neurips.cc/paper_files/paper/2024/hash/6ffdbbe354893979367f93e2121e37dd-Abstract-Conference.html
- Tang, J., Xu, J., Lu, T., Zhang, Z., et al. (2025). *Perception compressor: A training-free prompt compression framework in long context scenarios*. Findings of NAACL 2025.
<https://aclanthology.org/2025.findings-naacl.229/>
- Li, Z., Liu, Y., Su, Y., & Collier, N. (2025). *Prompt compression for large language models: A survey*. Proceedings of NAACL 2025.
<https://aclanthology.org/2025.naacl-long.368/>
- Faridi, A. R., Masood, F., Keshvi, et al. (2026). *Retrieval-Augmented Generation for Large Language Models: Evolution, Architectures, Applications, and Challenges (2020–2025)*. Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery.
<https://wires.onlinelibrary.wiley.com/doi/abs/10.1002/widm.70122>
- Guo, J., Su, Y., Ge, Y., Guan, S., Yu, L., et al. (2026). *Lost in Decomposition: Analyzing and Mitigating the Limitations of Long Context Methods via Context Dependency*. Findings of ACL 2026.
<https://aclanthology.org/2026.findings-acl.2097/>

Joshi, H., Shethia, P., Dao, J., & Lam, M. S. (2026). *Contexts are Never Long Enough: Structured Reasoning for Scalable Question Answering over Long Document Sets*. arXiv.

<https://arxiv.org/abs/2604.22294>

Ye, X., Zhang, W., Yin, F., Yen, H., & Chen, D. (2026). *DySCO: Dynamic Attention-Scaling Decoding for Long-Context Language Models*. arXiv.

<https://arxiv.org/abs/2602.22175>

Zhang, A. L., Kraska, T., & Khattab, O. (2025). *Recursive language models*. arXiv.

<https://arxiv.org/abs/2512.24601>